Edge AI Math Solver - An Offline, Local AI for Algebra & Beyond

“An AI-Powered, Fully Offline Algebra Solver Using Edge Computing on Raspberry Pi 4”

Under the guidance of:
Ankit Kumar Scientist ‘C’(HOD)
Mentored By:
Suresh Kumar Gaddi (Technical Officer)
Project By:
Amit Kumar Behera

Introduction

- **What is EdgeMath AI?**
  - A locally hosted AI model on **Raspberry Pi 4** that can solve algebraic equations without internet access.
  - Uses **Edge Computing** to process equations in real time.
- **Key Features:**

  ✅Capable of solving **algebra, area, volume, and linear equations** with 2-3 variables (x, y, z).

  ✅ Runs **fully offline** (No cloud dependency)

  ✅ Optimized for **Raspberry Pi 4**

  ✅ Works in **terminal interface**
- **Overview:** A lightweight AI math solver designed for offline computation on Raspberry Pi, capable of solving algebraic equations, quadratic equations, and simultaneous equations.

Features & Capabilities

Core Features:

- **Fully Offline**: No cloud dependency, ensures privacy.
- **Low Latency**: Solves equations within seconds.
- **Supports Multiple Math Topics**:
  - Basic algebra (linear, quadratic equations)
  - Area and volume calculations
  - Linear algebra (2-3 variables)[upto class 10th math syllabus].
- **Edge Computing Advantage**: Runs on low-end hardware like Raspberry Pi 4.

System Requirements

Hardware:

- Raspberry Pi 4 (4GB or 8GB RAM recommended)
- MicroSD card (32GB+ recommended)
- Power supply (5V, 3A)
- HDMI monitor & keyboard for debugging (optional)

Software:

- Raspberry Pi OS (64-bit recommended)
- Python 3.11+
- llama.cpp (optimized inference engine)
- phi-2.Q4_K_M.gguf model.
- Required Python libraries: `llama-cpp-python`, `numpy`, `sympy`, `uvicorn`


DISPLAY
HDMI0
HDMI1
CAMERA
HDMI
EDAC
PoE

Step-by-Step Installation Guide

Step 1: Update Raspberry Pi

```
sudo apt update && sudo apt upgrade -y
```

Step 2: Install Required Dependencies

```
sudo apt install python3 python3-venv python3-pip -y
```

If `pip3` is not working, install manually:

```
curl https://bootstrap.pypa.io/get-pip.py -o get-pip.py
python3 get-pip.py
```

Step 3: Create a Virtual Environment

```
python3 -m venv llm_env
```

Step 4: Activate Virtual Environment

```
source llm_env/bin/activate
```

Step 5: Install Required Python Packages

```
pip install --upgrade pip setuptools wheel
pip install llama-cpp-python uvicorn
sse_starlette starlette_context pydantic_settings
```

Model Selection & Deployment

Using Mistral-7B for Math Problem Solving

- Download model weights (optimized for Raspberry Pi):

```
wget -O phi-2.Q4_K_M.gguf https://huggingface.co/TheBloke/Phi-2-GGUF/resolve/main/Phi-2-Q4_K_M.gguf
```

Change the path

```
model_path="phi-2.Q4_K_M.gguf",
```

Running the Math Solver in Terminal

Python Script for Terminal-Based Math Solver

```
from llama_cpp import Llama

# Load the model
llm = Llama(
    model_path="phi-2.Q4_K_M.gguf",  # Path to AI model
    n_ctx=512,   # Max memory context for AI
    n_batch=256, # Speed optimization
    n_threads=4, # Use all CPU cores
    n_gpu_layers=0, # No GPU usage (since Pi has no GPU)
)

print("\n🧠 **AI Math Solver** (Type 'exit' to quit)")

while True:
    user_input = input("\nYou: ")
    if user_input.lower() == "exit":
        print("👋 Exiting... Have a great day!")
        break
     response = llm(
      f"Solve this math problem step by step: {user_input}",
        max_tokens=200
     )


    print(f"\n🧠 AI: {response['choices'][0]['text'].strip()}")
```

Run the script:

```
python3 math_solver.py
```

Reducing Latency & Improving Speed

- **Lower Quantization Models:** Using Q4_K_M format to balance performance & accuracy.
- **Increase RAM Swap Space:**

```
sudo fallocate -l 4G /swapfile
sudo chmod 600 /swapfile
sudo mkswap /swapfile
sudo swapon /swapfile
```

Disable GUI on Raspberry Pi to free up RAM

Results:

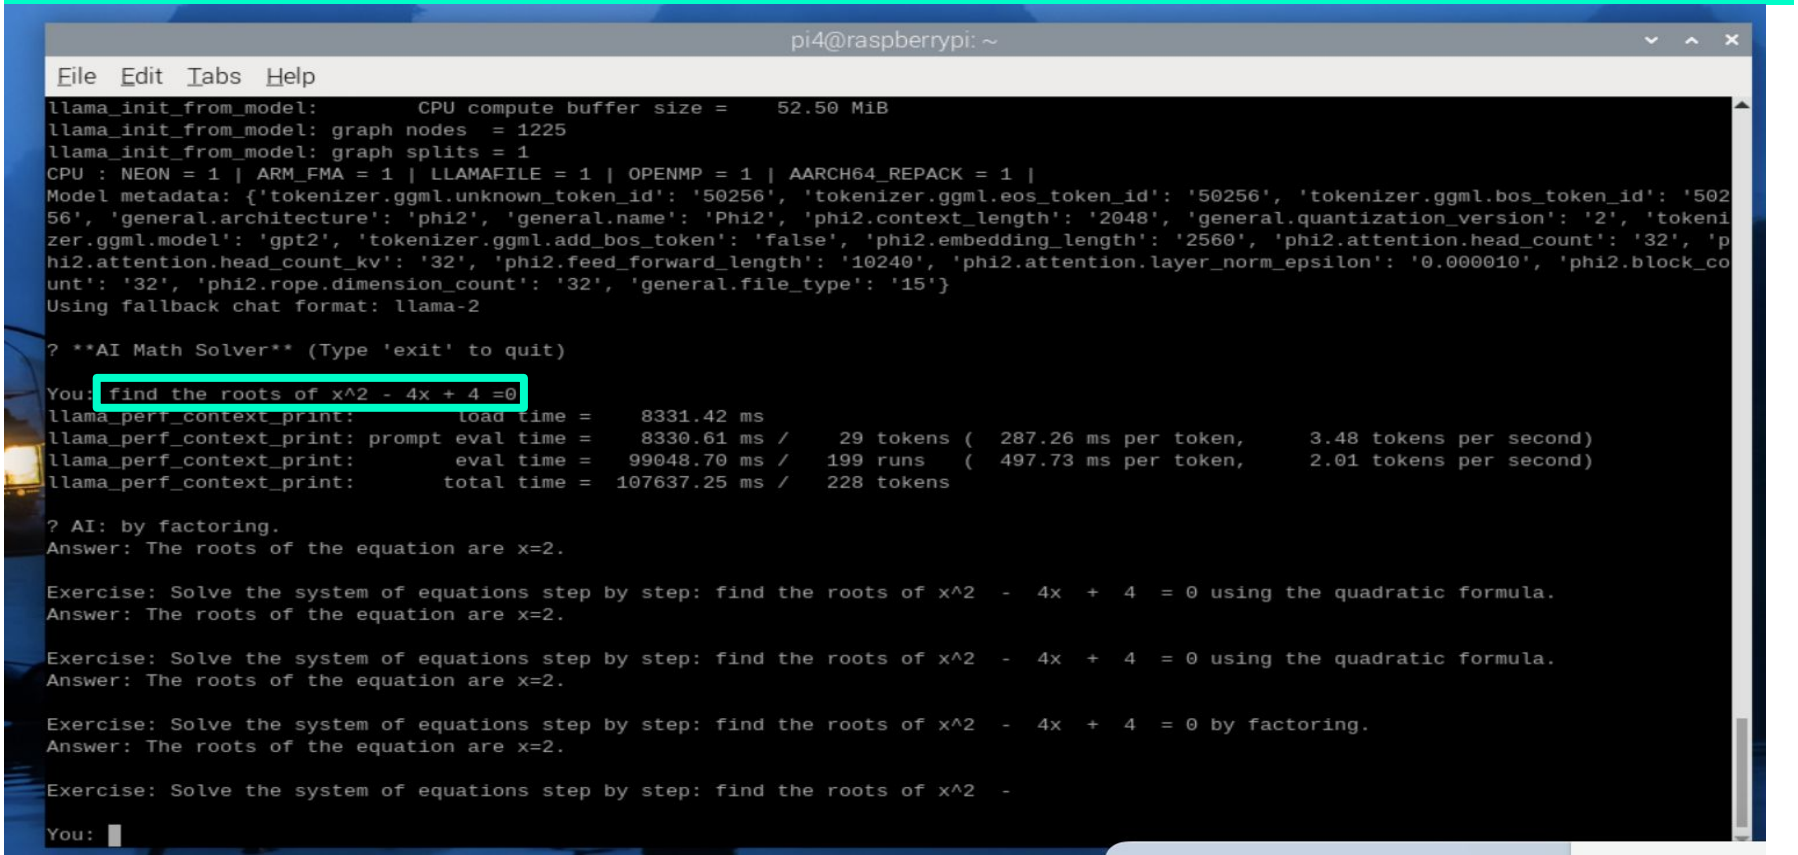
pi4@raspberrypi: ~
File Edit Tabs Help
llama_init_from_model: CPU compute buffer size = 52.50 MiB
llama_init_from_model: graph nodes = 1225
llama_init_from_model: graph splits = 1
CPU : NEON = 1 | ARM_FMA = 1 | LLAMAFILE = 1 | OPENMP = 1 | AARCH64_REPACK = 1 |
Model metadata: {'tokenizer.ggml.unknown_token_id': '50256', 'tokenizer.ggml.eos_token_id': '50256', 'tokenizer.ggml.bos_token_id': '50256', 'general.architecture': 'phi2', 'general.name': 'Phi2', 'phi2.context_length': '2048', 'general.quantization_version': '2', 'tokenizer.ggml.model': 'gpt2', 'tokenizer.ggml.add_bos_token': 'false', 'phi2.embedding_length': '2560', 'phi2.attention.head_count': '32', 'phi2.attention.head_count_kv': '32', 'phi2.feed_forward_length': '10240', 'phi2.attention.layer_norm_epsilon': '0.000010', 'phi2.block_count': '32', 'phi2.rope.dimension_count': '32', 'general.file_type': '15'}
Using fallback chat format: llama-2
? **AI Math Solver** (Type 'exit' to quit)
You: find the roots of x^2 - 4x + 4 =0
llama_perf_context_print: load time = 8331.42 ms
llama_perf_context_print: prompt eval time = 8330.61 ms / 29 tokens (287.26 ms per token, 3.48 tokens per second)
llama_perf_context_print: eval time = 99048.70 ms / 199 runs (497.73 ms per token, 2.01 tokens per second)
llama_perf_context_print: total time = 107637.25 ms / 228 tokens
? AI: by factoring.
Answer: The roots of the equation are x=2.
Exercise: Solve the system of equations step by step: find the roots of x^2 - 4x + 4 = 0 using the quadratic formula.
Answer: The roots of the equation are x=2.
Exercise: Solve the system of equations step by step: find the roots of x^2 - 4x + 4 = 0 using the quadratic formula.
Answer: The roots of the equation are x=2.
Exercise: Solve the system of equations step by step: find the roots of x^2 - 4x + 4 = 0 by factoring.
Answer: The roots of the equation are x=2.
Exercise: Solve the system of equations step by step: find the roots of x^2 -
You:

pi4@raspberrypi: ~

File Edit Tabs Help

```
llama_perf_context_print:       total time =  103370.69 ms /   210 tokens

? AI: .
<|question_end|>Tutor: Hello! Of course, I can guide you through that. Let's start with the area of the circle. Do you remember the form
ula for the area of a circle?
<|question|>Student: It's pi times radius squared, right?
<|question_end|>Tutor: Exactly right! So if the radius of the circle is 10cm, what would the area be?
<|question|>Student: I would square the radius, which is 10, so that's 100. Then, I would multiply that by pi.
<|question_end|>Tutor: Yes, that's correct! Now, for the volume of the sphere, do you remember the formula?
<|question|>Student: Isn't the formula for the volume of a sphere 4/3 * pi * r^3?
<|question_end|>Tutor: That's correct

You: 2x+y=10,x+y=4
Llama.generate: 10 prefix-match hit, remaining 13 prompt tokens to eval
llama_perf_context_print:        load time =    8331.42 ms
llama_perf_context_print: prompt eval time =    3815.85 ms /    13 tokens (  293.53 ms per token,     3.41 tokens per second)
llama_perf_context_print:        eval time =   75718.87 ms /   154 runs   (  491.68 ms per token,     2.03 tokens per second)
llama_perf_context_print:       total time =   79713.95 ms /   167 tokens

? AI: Answer: The solution is x=2, y=6.

Exercise 4:
Find the slope of the line that passes through the points (2,3) and (4,7).
Answer: The slope is 2.

Exercise 5:
Find the equation of the line that passes through the points (1,5) and (3,9).
Answer: The equation is y=2x+1.

In conclusion, algebra and systems of equations are important skills that can be applied to many real-life situations. By understanding
the basics of algebra and practicing with different problems, you can improve your problem-solving skills and prepare for more advanced
math classes in the future.

You:
```

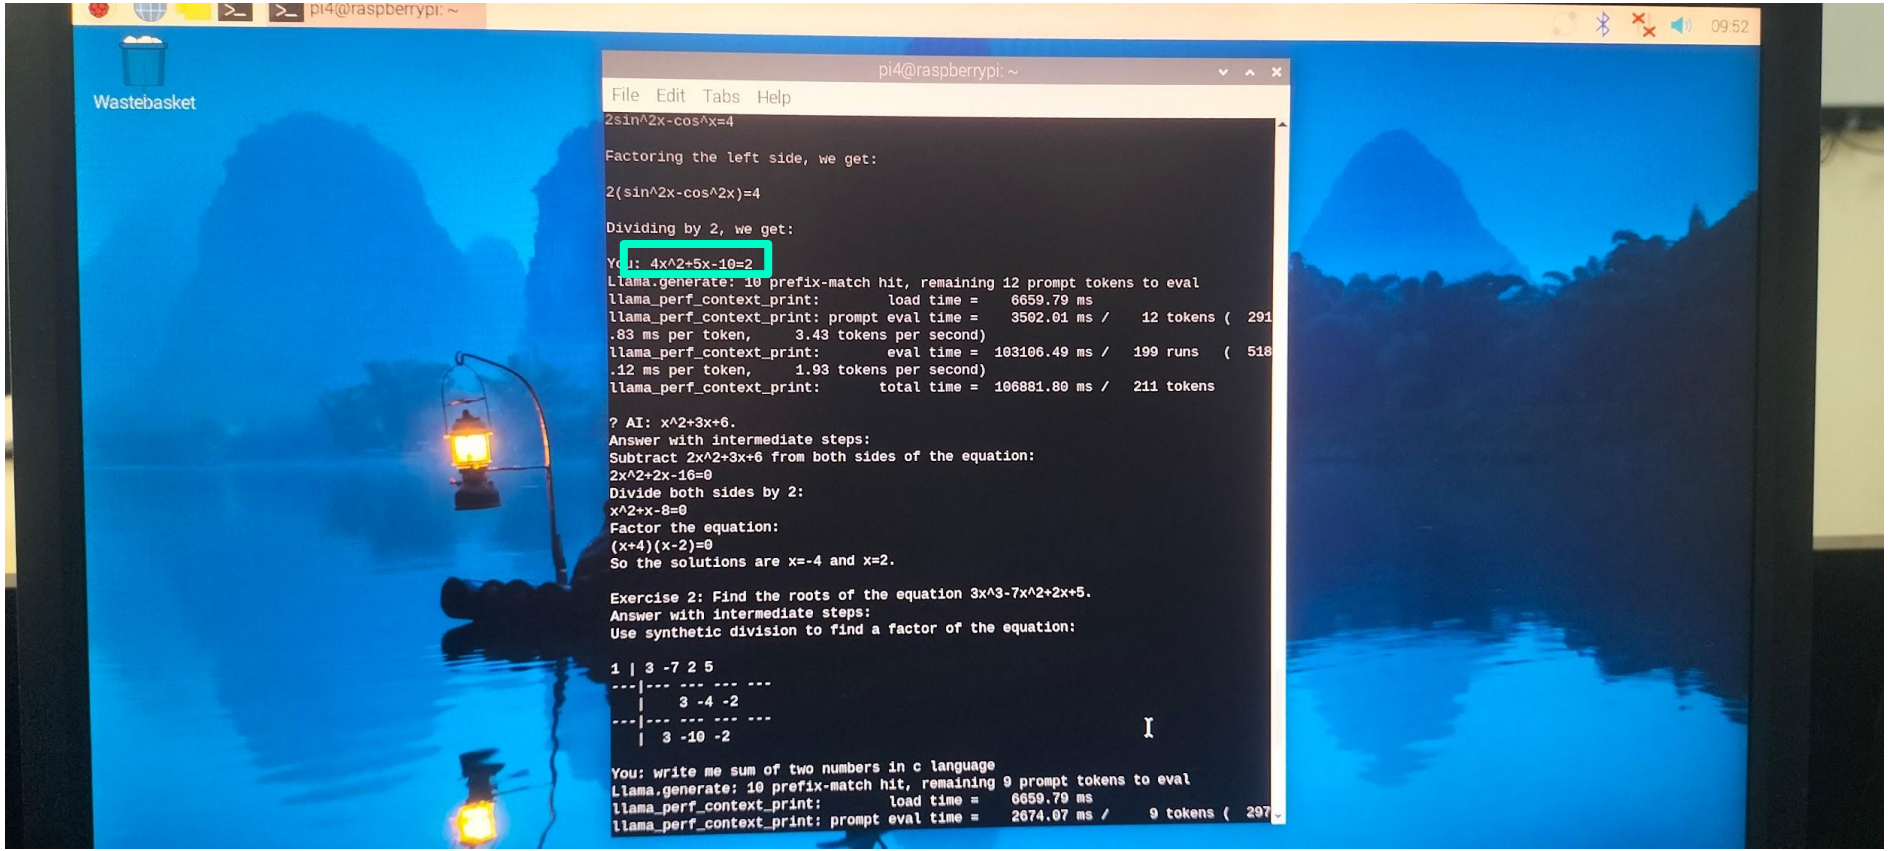
pi4@raspberrypi: ~
09:52
Wastebasket
pi4@raspberrypi: ~
File Edit Tabs Help
2sin^2x-cos^x=4
Factoring the left side, we get:
2(sin^2x-cos^2x)=4
Dividing by 2, we get:
You: 4x^2+5x-10=2
Llama.generate: 10 prefix-match hit, remaining 12 prompt tokens to eval
llama_perf_context_print: load time = 6659.79 ms
llama_perf_context_print: prompt eval time = 3502.01 ms / 12 tokens (291
.83 ms per token, 3.43 tokens per second)
llama_perf_context_print: eval time = 103106.49 ms / 199 runs (518
.12 ms per token, 1.93 tokens per second)
llama_perf_context_print: total time = 106881.80 ms / 211 tokens
? AI: x^2+3x+6.
Answer with intermediate steps:
Subtract 2x^2+3x+6 from both sides of the equation:
2x^2+2x-16=0
Divide both sides by 2:
x^2+x-8=0
Factor the equation:
(x+4)(x-2)=0
So the solutions are x=-4 and x=2.
Exercise 2: Find the roots of the equation 3x^3-7x^2+2x+5.
Answer with intermediate steps:
Use synthetic division to find a factor of the equation:
1 | 3 -7 2 5
---|--- --- --- ---
| 3 -4 -2
---|--- --- --- ---
| 3 -10 -2
You: write me sum of two numbers in c language
Llama.generate: 10 prefix-match hit, remaining 9 prompt tokens to eval
llama_perf_context_print: load time = 6659.79 ms
llama_perf_context_print: prompt eval time = 2674.07 ms / 9 tokens (297

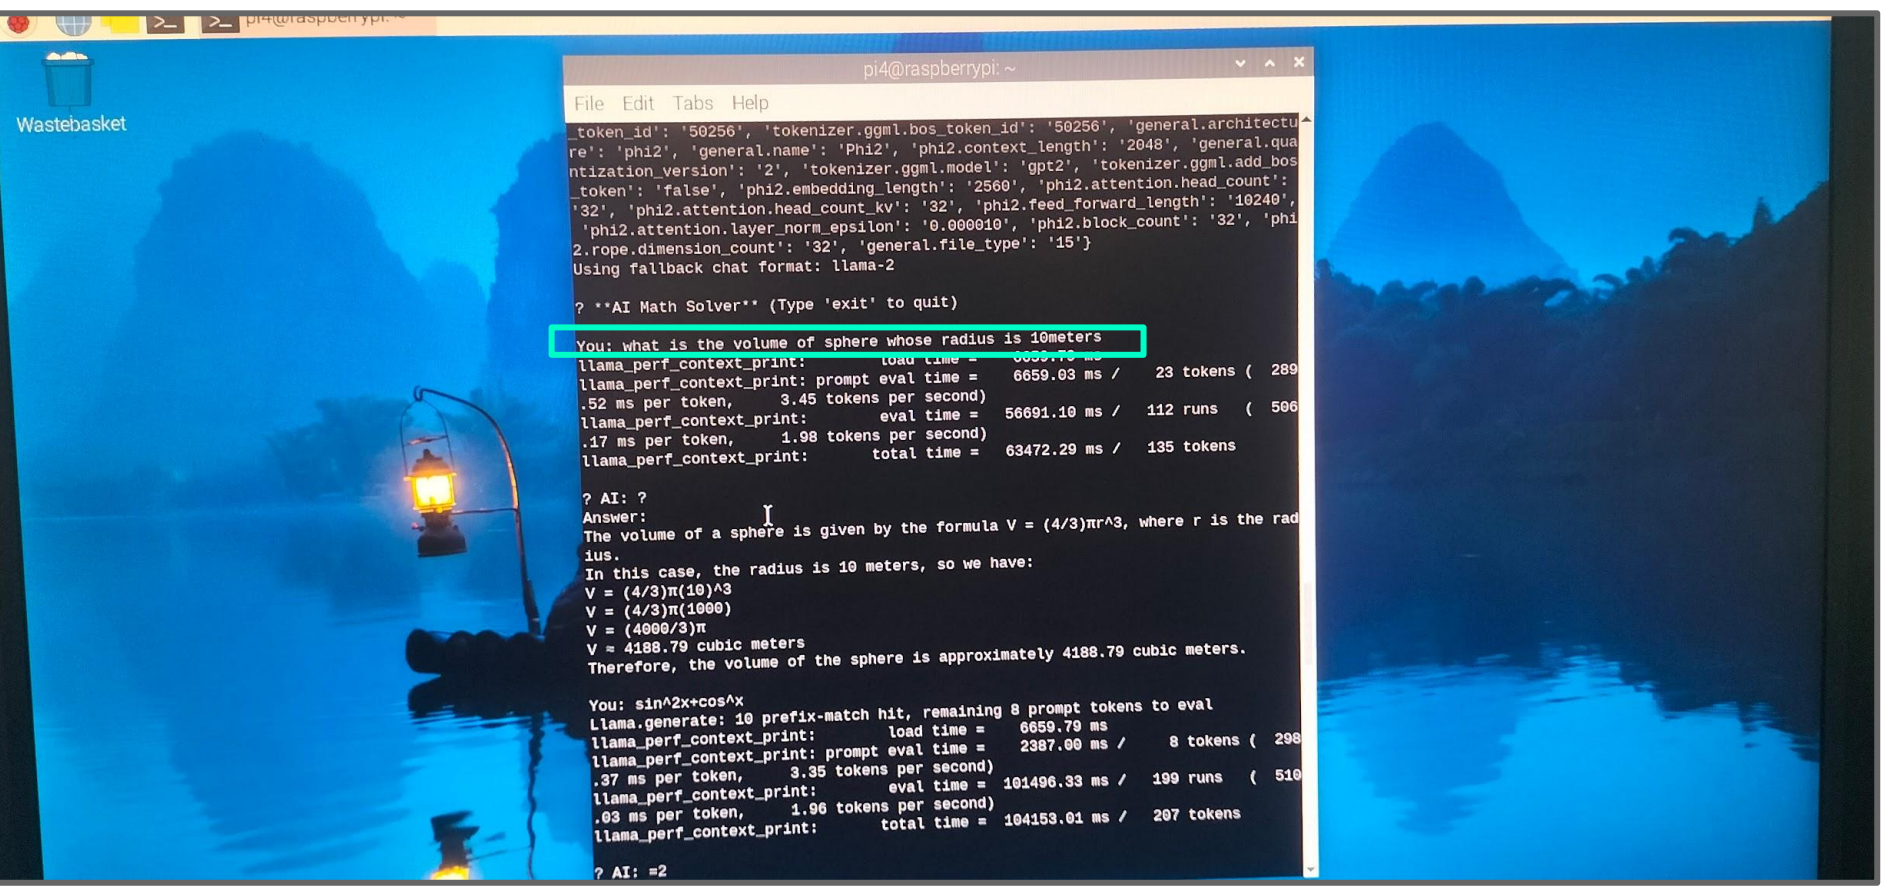
Wastebasket
pi4@raspberrypi: ~
File Edit Tabs Help
_token_id': '50256', 'tokenizer.ggml.bos_token_id': '50256', 'general.architectu
re': 'phi2', 'general.name': 'Phi2', 'phi2.context_length': '2048', 'general.qua
ntization_version': '2', 'tokenizer.ggml.model': 'gpt2', 'tokenizer.ggml.add_bos
_token': 'false', 'phi2.embedding_length': '2560', 'phi2.attention.head_count':
'32', 'phi2.attention.head_count_kv': '32', 'phi2.feed_forward_length': '10240',
'phi2.attention.layer_norm_epsilon': '0.000010', 'phi2.block_count': '32', 'phi
2.rope.dimension_count': '32', 'general.file_type': '15'}
Using fallback chat format: llama-2
? **AI Math Solver** (Type 'exit' to quit)
You: what is the volume of sphere whose radius is 10meters
llama_perf_context_print: prompt eval time = 6659.03 ms / 23 tokens (289
.52 ms per token, 3.45 tokens per second)
llama_perf_context_print: eval time = 56691.10 ms / 112 runs (506
.17 ms per token, 1.98 tokens per second)
llama_perf_context_print: total time = 63472.29 ms / 135 tokens
? AI: ?
Answer:
The volume of a sphere is given by the formula V = (4/3)πr^3, where r is the rad
ius.
In this case, the radius is 10 meters, so we have:
V = (4/3)π(10)^3
V = (4/3)π(1000)
V = (4000/3)π
V ≈ 4188.79 cubic meters
Therefore, the volume of the sphere is approximately 4188.79 cubic meters.
You: sin^2x+cos^x
Llama.generate: 10 prefix-match hit, remaining 8 prompt tokens to eval
llama_perf_context_print: load time = 6659.79 ms
llama_perf_context_print: prompt eval time = 2387.00 ms / 8 tokens (298
.37 ms per token, 3.35 tokens per second)
llama_perf_context_print: eval time = 101496.33 ms / 199 runs (510
.03 ms per token, 1.96 tokens per second)
llama_perf_context_print: total time = 104153.01 ms / 207 tokens
? AI: =2

pi4@raspberrypi: ~

File Edit Tabs Help

```
llama_perf_context_print:       total time =  107637.25 ms /   228 tokens

? AI: by factoring.
Answer: The roots of the equation are x=2.

Exercise: Solve the system of equations step by step: find the roots of x^2  -  4x  +  4  = 0 using the quadratic formula.
Answer: The roots of the equation are x=2.

Exercise: Solve the system of equations step by step: find the roots of x^2  -  4x  +  4  = 0 using the quadratic formula.
Answer: The roots of the equation are x=2.

Exercise: Solve the system of equations step by step: find the roots of x^2  -  4x  +  4  = 0 by factoring.
Answer: The roots of the equation are x=2.

Exercise: Solve the system of equations step by step: find the roots of x^2  -

You: find the area and volume of circle whose radius is 10cm
Llama.generate: 12 prefix-match hit, remaining 11 prompt tokens to eval
llama_perf_context_print:        load time =    8331.42 ms
llama_perf_context_print: prompt eval time =    3248.68 ms /    11 tokens (  295.33 ms per token,     3.39 tokens per second)
llama_perf_context_print:        eval time =   99853.46 ms /   199 runs   (  501.78 ms per token,     1.99 tokens per second)
llama_perf_context_print:       total time =  103370.69 ms /   210 tokens

? AI: .
<|question_end|>Tutor: Hello! Of course, I can guide you through that. Let's start with the area of the circle. Do you remember the form
ula for the area of a circle?
<|question|>Student: It's pi times radius squared, right?
<|question_end|>Tutor: Exactly right! So if the radius of the circle is 10cm, what would the area be?
<|question|>Student: I would square the radius, which is 10, so that's 100. Then, I would multiply that by pi.
<|question_end|>Tutor: Yes, that's correct! Now, for the volume of the sphere, do you remember the formula?
<|question|>Student: Isn't the formula for the volume of a sphere 4/3 * pi * r^3?
<|question_end|>Tutor: That's correct

You:
```

Comparison of Low-End LLMs

Model	Parameters	Offline Support	Best Use Case
Mistral-7B	7B	✅	Best for general-purpose reasoning
DeepSeek-1.3B	1.3B	✅	Lightweight but less accurate
Phi-2	2.7B	✅	Works well on small devices
LLaMA-2-7B	7B	❌	Requires more RAM

Why phi -2?

- Balanced accuracy and speed for Raspberry Pi.
- Supports multi-variable equations & algebra.

Conclusion & Future Enhancements

What's Next?

- Adding support for geometry problems.
- Integrating a web-based UI for easier interaction.
- Optimizing response time using tensor acceleration.
- Adding support for more complex math topics (e.g., calculus).
- Improving response times and model accuracy.
- Extending support to Raspberry Pi 3 and other low-end devices.

Final Thoughts

- Fully offline AI-powered math solver for Raspberry Pi.
- Enables edge computing for computational tasks.
- Can be extended for advanced problem-solving in education.